

# Foundational Machine Learning

## Week 3 : Random Forests and Bias Variance Tradeoff

Rohit Babbar  
rb2608@bath.ac.uk



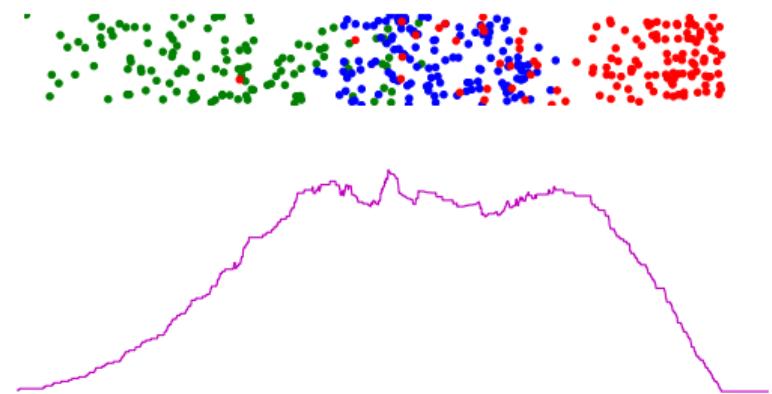
## This lecture

- Last week:
  - Decision trees
    - Input = real
    - Output = categorical (discrete)

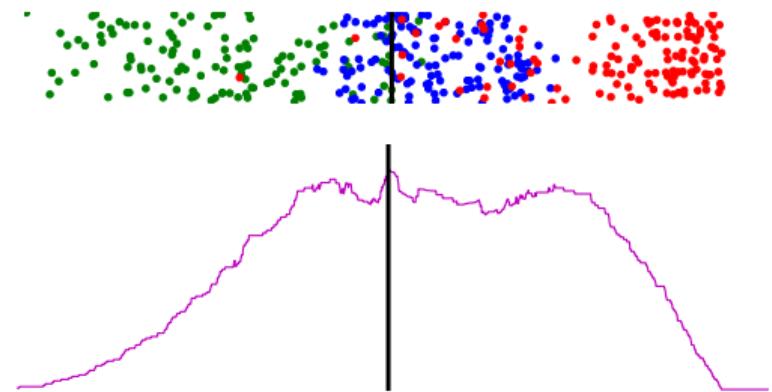
- Last week:
  - Decision trees
    - Input = real
    - Output = categorical (discrete)
- This week:
  - Decision trees with some further input and output types
  - Bias-variance tradeoff – some theoretical insights
  - Bagging - an ensembling technique
  - **Random forests**

# Inputs

- Choose best Gini impurity / info gain for all axes  
(1 axis shown; vertical offset for visualisation only)

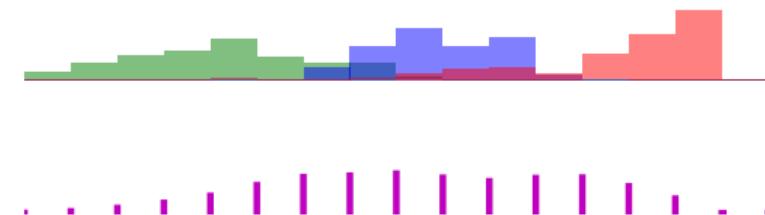


- Choose best Gini impurity / info gain for all axes  
(1 axis shown; vertical offset for visualisation only)



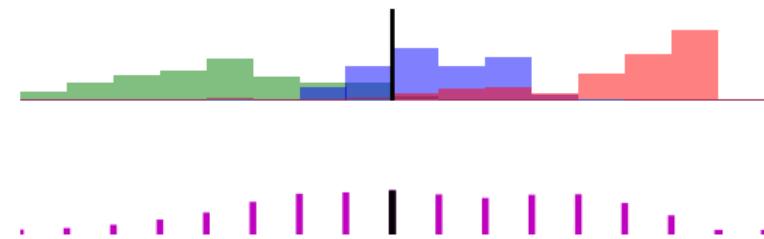
## Quantised real input

- Continuous data may be **quantised**
- e.g. “What is your age?”
  - From 12–17, 18–24 etc.
  - Or year only



## Quantised real input

- Continuous data may be **quantised**
- e.g. “What is your age?”
  - From 12–17, 18–24 etc.
  - Or year only
- Split between bins
  - (Information gain shown as spikes as only defined at bin transitions)



## Categorical input

- Similar to quantised (histogram of categories) . . .  
    . . . but unordered
- e.g. "What is your favourite cheese?"
- Splitting no longer makes sense!



## Categorical input

- Similar to quantised (histogram of categories) . . .  
    . . . but unordered
- e.g. "What is your favourite cheese?"
- Splitting no longer makes sense!
- Two choices :
  - Try every assignment of category to the left/right side and pick best  
(one category always goes left to account for symmetry)
  - One category goes down left branch, rest go right



## Categorical input

- Similar to quantised (histogram of categories) . . .  
    . . . but unordered
- e.g. "What is your favourite cheese?"
- Splitting no longer makes sense!
- Two choices :
  - Try every assignment of category to the left/right side and pick best  
(one category always goes left to account for symmetry)
  - One category goes down left branch, rest go right
- One left, rest right is preferred:
  - Fixed storage
  - Simpler code
  - Combinations to test grows linearly  
( $\mathcal{O}(n)$ , not  $\mathcal{O}(2^{n-1})$ ); where  $n$  = numbers of categories)



# Outputs

- Classification:
  - Split to minimise Gini impurity or maximise information gain
  - Leaf gives answer as most common class to reach it

- Classification:
  - Split to minimise Gini impurity or maximise information gain
  - Leaf gives answer as most common class to reach it
- Regression:
  - Split to minimise variance or maximise information gain
  - Leaf gives answer as mean value to reach it
    - (median may confer an advantage — any idea when?)
- Otherwise identical!

## Variance reduction

- Variance of output measures how consistent a node is...  
...so choose splits that minimise it

- Variance of output measures how consistent a node is...  
...so choose splits that minimise it
- Variance of left node:

$$\sigma_l^2 = \mathbb{E} [(Y_l - \mathbb{E}[Y_l])^2]$$

similarly for right,  $\sigma_r^2$  ( $Y_l$  = data that goes left)

- Variance of output measures how consistent a node is...  
...so choose splits that minimise it
- Variance of left node:

$$\sigma_l^2 = \mathbb{E} [(Y_l - \mathbb{E}[Y_l])^2]$$

similarly for right,  $\sigma_r^2$  ( $Y_l$  = data that goes left)

- Minimise weighted combination:

$$L(\text{split}) = \frac{n_l}{n} \sigma_l^2 + \frac{n_r}{n} \sigma_r^2$$

$n$  = total exemplar count

$n_l$  = exemplars traveling down left branch

$n_r$  = exemplars traveling down right branch

## Information gain

- Information gain is not often used for regression tasks with decision trees, as it is not immediately applicable

---

<sup>1</sup><https://www.biopsychology.org/norwich/isp/chap8.pdf>

- Information gain is not often used for regression tasks with decision trees, as it is not immediately applicable

In order to apply,

- First fit Gaussian distribution to output variable
- Then, compute entropy <sup>1</sup>

$$\frac{1}{2} \log (2\pi e \sigma^2)$$

- Information gain, thus, is

$$I(\text{split}) = \frac{1}{2} \log (2\pi e \sigma_p^2) - \frac{n_l}{2n} \log (2\pi e \sigma_l^2) - \frac{n_r}{2n} \log (2\pi e \sigma_r^2)$$

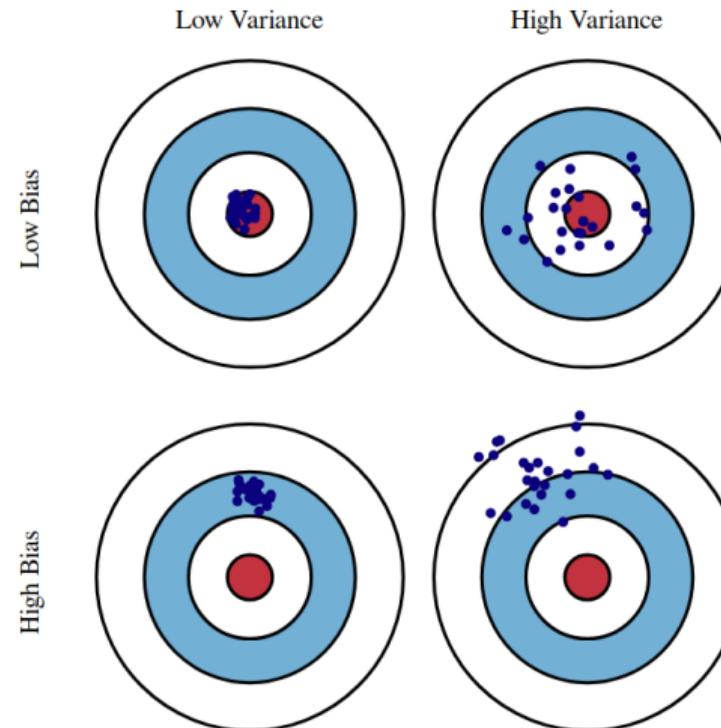
(same variables as previous slide, with  $p$  subscript for parent)

---

<sup>1</sup><https://www.biopsychology.org/norwich/isp/chap8.pdf>

From Decision Trees to Random Forests via Bias-variance tradeoff

## Bias-Variance Tradeoff



**Figure:** Pictorial depiction of the components of bias-variance tradeoff

## Unbiased Estimator Concept

### Definition

An estimator  $\hat{\theta}$  of a population parameter  $\theta$  is **unbiased** if:

$$\mathbb{E}[\hat{\theta}] = \theta$$

That is, on average, it neither overestimates nor underestimates  $\theta$ .

## Sample Mean Setup

Let  $X_1, X_2, \dots, X_n$  be i.i.d. random variables with:

$$\mathbb{E}[X_i] = \mu \quad \text{and} \quad \text{Var}(X_i) = \sigma^2$$

The sample mean is defined as:

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

## Expectation of the Sample Mean

Using the linearity of expectation:

$$\mathbb{E}[\bar{X}] = \mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n X_i\right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}[X_i] = \frac{1}{n} \cdot n\mu = \mu$$

### Conclusion

$$\boxed{\mathbb{E}[\bar{X}] = \mu}$$

The sample mean is therefore an **unbiased estimator** of the population mean.

## A common way to compute Sample variance

Biased estimator (uses  $n$  in denominator)

$$\tilde{s}^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2$$

**Key identity:**

$$\sum_{i=1}^n (X_i - \bar{X})^2 = \sum_{i=1}^n (X_i - \mu)^2 - n(\bar{X} - \mu)^2$$

**Take expectations (i.i.d.,  $E[X_i] = \mu$ ,  $\text{Var}(X_i) = \sigma^2$ ):**

$$\mathbb{E}\left[\sum_{i=1}^n (X_i - \mu)^2\right] = n\sigma^2, \quad \mathbb{E}\left[n(\bar{X} - \mu)^2\right] = n \cdot \text{Var}(\bar{X}) = n \cdot \frac{\sigma^2}{n} = \sigma^2$$

**Therefore:**

$$\mathbb{E}\left[\tilde{s}^2\right] = \frac{1}{n} \left(n\sigma^2 - \sigma^2\right) = \left(1 - \frac{1}{n}\right) \sigma^2 = \frac{n-1}{n} \sigma^2$$

Turns out that using  $(n-1)$  in the denominator makes it unbiased

## Bias-variance tradeoff<sup>2 3</sup>

Consider regression problem with squared error

- Given a training set  $D = \{(x_i, y_i)\}_{i=1}^n$  such that  $x_i \in \mathbb{R}^d$  and  $y_i \in \mathbb{R}$ .
- Assume that  $y_i = f(x_i) + \epsilon$ , where  $\epsilon$  is a random variable representing noise with mean 0 and variance  $\sigma^2$
- Since  $f(\cdot)$  is unknown, we try to approximate it using the training data, and let  $\hat{f}(\cdot)$  denote our approximation. Imagine  $\hat{f}(\cdot)$  to be decision tree (for regression) that we constructed,
- In regression, this is done by minimising the squared error between  $y$  and  $f(x)$ , i.e.  $(y - \hat{f}(x))^2$ , where the pair  $(x, y)$  could be a training data point or a novel/unseen test data point.

---

<sup>2</sup>Based on Wikipedia article [https://en.wikipedia.org/wiki/Bias-variance\\_tradeoff](https://en.wikipedia.org/wiki/Bias-variance_tradeoff)

<sup>3</sup>Derivation as a whole is non-examinable, but you should still have an understanding of the individual parts and the overall idea of the proof

## Bias-variance tradeoff - II

Irrespective of how the classifier  $\hat{f}(\cdot)$  is learnt on the data, it's **expected** error on an unseen sample (test) sample  $x$  can be decomposed as follows :

$$\mathbb{E}_{D,\varepsilon} \left[ (y - \hat{f}(x; D))^2 \right] = \left( \text{Bias}_D [\hat{f}(x; D)] \right)^2 + \text{Var}_D [\hat{f}(x; D)] + \sigma^2$$

where

$$\text{Bias}_D [\hat{f}(x; D)] \triangleq \mathbb{E}_D [\hat{f}(x; D)] - f(x)$$

$$\text{Var}_D [\hat{f}(x; D)] \triangleq \mathbb{E}_D \left[ (\mathbb{E}_D [\hat{f}(x; D)] - \hat{f}(x; D))^2 \right]$$

$$\sigma^2 = \mathbb{E}_y \left[ (y - f(x))^2 \right]$$

Note  $\triangleq$  means this is a definition

## Bias-variance tradeoff - III

- Rewriting the LHS of the equation on the previous slide as MSE (Mean squared error)

$$\begin{aligned}
 \text{MSE} &\triangleq \mathbb{E}\left[\left(y - \hat{f}(x)\right)^2\right] && \text{writing the expectation } \mathbb{E} \text{ without subscripts} \\
 &= \mathbb{E}\left[\left(f(x) + \varepsilon - \hat{f}(x)\right)^2\right] && \text{since } y \triangleq f(x) + \varepsilon \\
 &= \mathbb{E}\left[\left(f(x) - \hat{f}(x)\right)^2\right] + 2 \mathbb{E}\left[\left(f(x) - \hat{f}(x)\right)\varepsilon\right] + \mathbb{E}[\varepsilon^2]
 \end{aligned}$$

- Using the fact that for independent r.v.  $X$  and  $Y$ ,  $\mathbb{E}(XY) = \mathbb{E}(X)\mathbb{E}(Y)$

$$\begin{aligned}
 \mathbb{E}\left[\left(f(x) - \hat{f}(x)\right)\varepsilon\right] &= \mathbb{E}[f(x) - \hat{f}(x)] \mathbb{E}[\varepsilon] && \text{since } \varepsilon \text{ is independent from } x \\
 &= 0 && \text{since } \mathbb{E}[\varepsilon] = 0
 \end{aligned}$$

- Expanding the first term below :

$$\begin{aligned}
 \mathbb{E}\left[\left(f(x) - \hat{f}(x)\right)^2\right] &= \mathbb{E}\left[\left(f(x) - \mathbb{E}[\hat{f}(x)] + \mathbb{E}[\hat{f}(x)] - \hat{f}(x)\right)^2\right] \\
 &= \mathbb{E}\left[\left(f(x) - \mathbb{E}[\hat{f}(x)]\right)^2\right] + 2 \mathbb{E}\left[\left(f(x) - \mathbb{E}[\hat{f}(x)]\right)\left(\mathbb{E}[\hat{f}(x)] - \hat{f}(x)\right)\right] \\
 &\quad + \mathbb{E}\left[\left(\mathbb{E}[\hat{f}(x)] - \hat{f}(x)\right)^2\right]
 \end{aligned}$$

## Bias-variance tradeoff - IV

$$\begin{aligned}
 \mathbb{E}[(f(x) - \mathbb{E}[\hat{f}(x)])^2] &= \mathbb{E}[f(x)^2] - 2 \mathbb{E}[f(x) \mathbb{E}[\hat{f}(x)]] + \mathbb{E}[\mathbb{E}[\hat{f}(x)]^2] \\
 &= f(x)^2 - 2 f(x) \mathbb{E}[\hat{f}(x)] + \mathbb{E}[\hat{f}(x)]^2 \\
 &= (f(x) - \mathbb{E}[\hat{f}(x)])^2
 \end{aligned}$$

The term in red :

$$\begin{aligned}
 \mathbb{E}[(f(x) - \mathbb{E}[\hat{f}(x)]) (\mathbb{E}[\hat{f}(x)] - \hat{f}(x))] &= \mathbb{E}[f(x) \mathbb{E}[\hat{f}(x)] - f(x) \hat{f}(x) - \mathbb{E}[\hat{f}(x)]^2 + \mathbb{E}[\hat{f}(x)] \hat{f}(x)] \\
 &= f(x) \mathbb{E}[\hat{f}(x)] - f(x) \mathbb{E}[\hat{f}(x)] - \mathbb{E}[\hat{f}(x)]^2 + \mathbb{E}[\hat{f}(x)]^2 \\
 &= 0
 \end{aligned}$$

$$\begin{aligned}
 \text{MSE} &= (f(x) - \mathbb{E}[\hat{f}(x)])^2 + \mathbb{E}[(\mathbb{E}[\hat{f}(x)] - \hat{f}(x))^2] + \sigma^2 \\
 &= \text{Bias}(\hat{f}(x))^2 + \text{Var}[\hat{f}(x)] + \sigma^2
 \end{aligned}$$

## Bias-variance tradeoff - V

- Variance: Captures how much the classifier changes if trained on a (slightly) different training set. How "over-specialized" is it to a particular training set (overfitting)?
- Bias: What is the inherent error that classifiers incurs even with infinite training data? This is due to the classifier being "biased" to a particular kind of solution (e.g. linear classifier).
- Noise: How much is the data-intrinsic noise? It's a measure of the ambiguity due to the data distribution and feature representation ( $\epsilon$  in the above example).
- The goal in Random Forest (next) is to lower the variance  $\mathbb{E}\left[\left(\mathbb{E}[\hat{f}(x)] - \hat{f}(x)\right)^2\right]$  by having more and more estimates of  $\hat{f}(x)$ , i.e. growing lots of decision trees into (random) forests !

# Random Forests

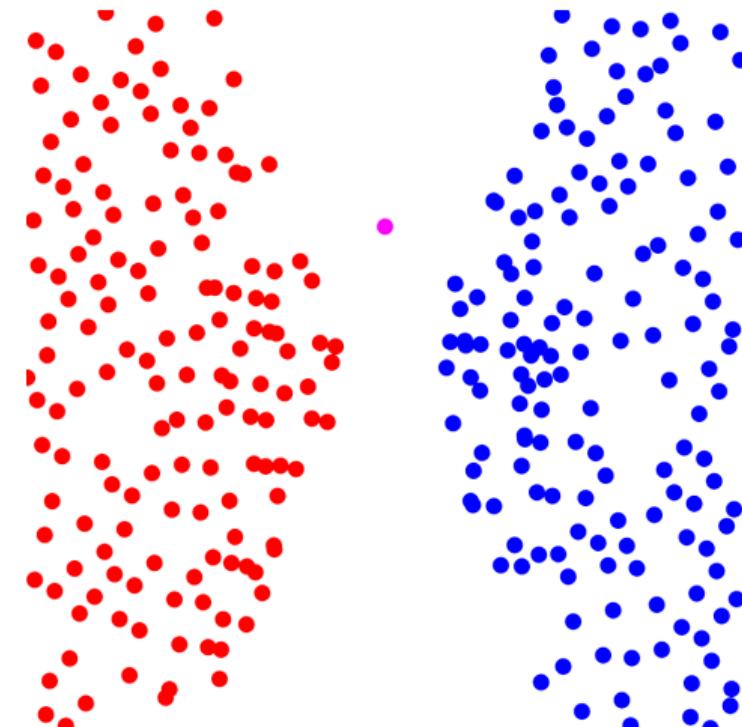
- Random forest = decision trees + ensemble learning

- Random forest = decision trees + ensemble learning
- Ensemble learning = combining multiple estimators
  - Different models (e.g. linear regression, decision tree, SVM, neural network)  
or
  - Same model, randomised training so each is different  
(using estimator to distinguish from model)

- Random forest = decision trees + ensemble learning
- Ensemble learning = combining multiple estimators
  - Different models (e.g. linear regression, decision tree, SVM, neural network)  
or
  - Same model, randomised training so each is different  
(using estimator to distinguish from model)
- Random forest:
  - Many decision trees (hence name)
  - Randomised training using **bagging**,  
a specific ensemble learning technique

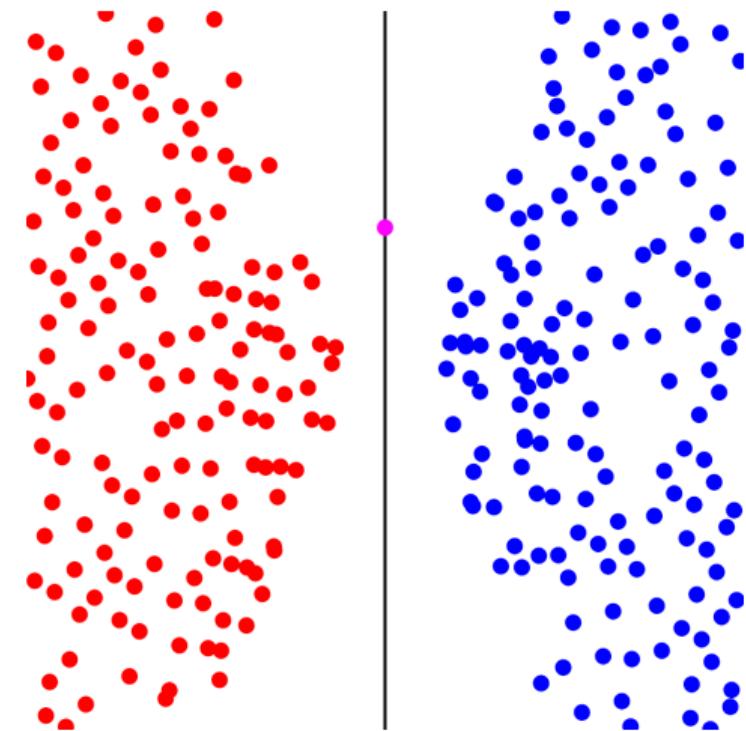
## Ensemble learning II

- Which class (red or blue) should the magenta dot be?



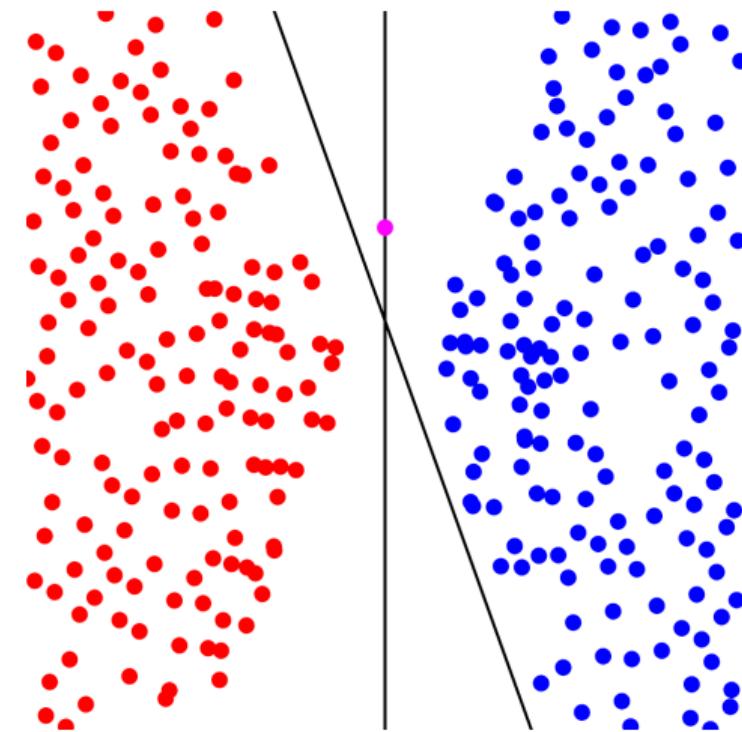
## Ensemble learning II

- Which class (red or blue) should the magenta dot be?
- An obvious classification boundary ...



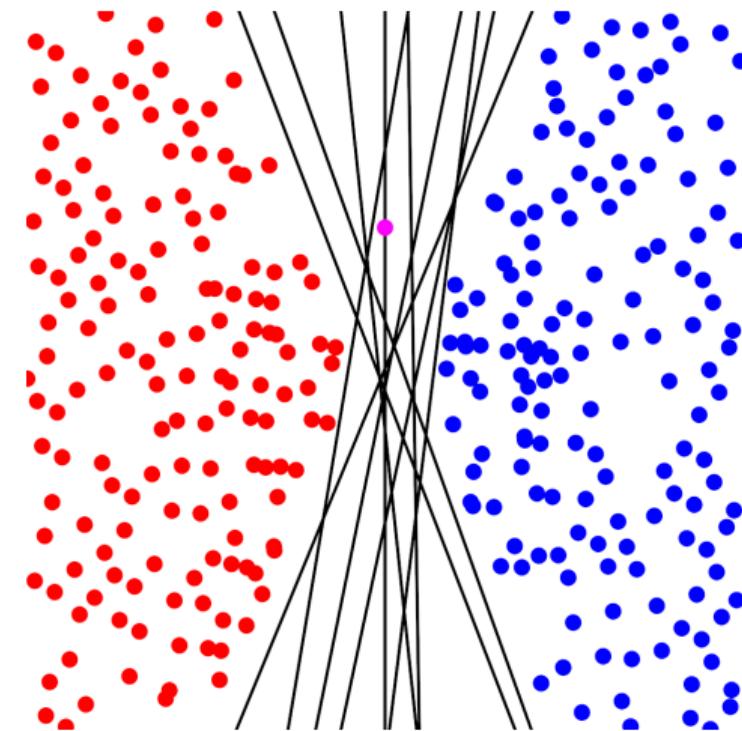
## Ensemble learning II

- Which class (red or blue) should the magenta dot be?
- An obvious classification boundary ...
- But this is just as good (suggesting blue)



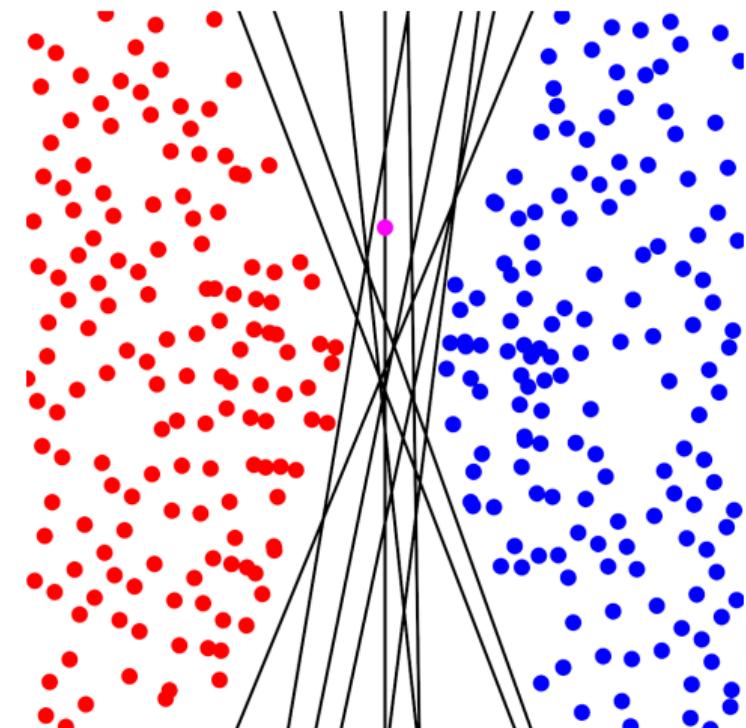
## Ensemble learning II

- Which class (red or blue) should the magenta dot be?
- An obvious classification boundary ...
- But this is just as good (suggesting blue)
- As are all of these!



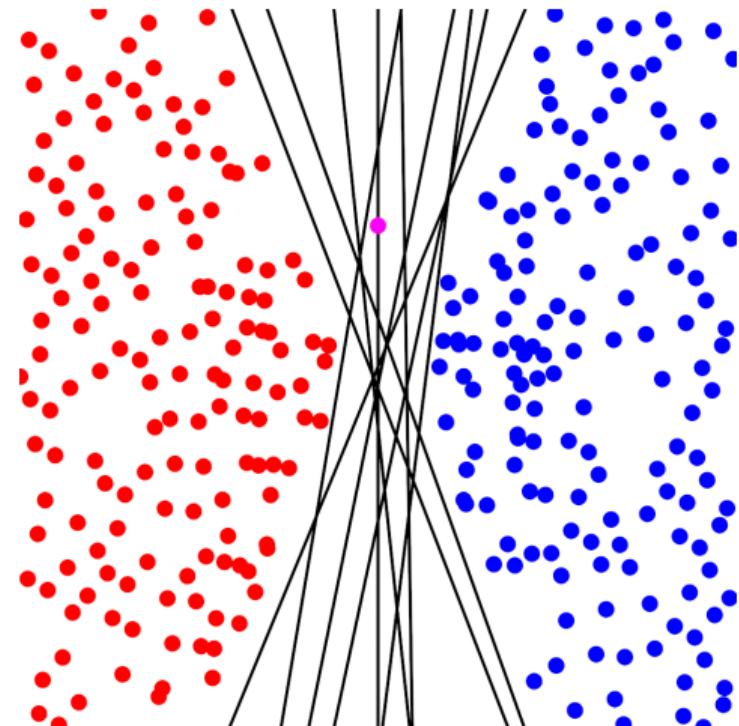
## Ensemble learning II

- Which class (red or blue) should the magenta dot be?
- An obvious classification boundary ...
- But this is just as good (suggesting blue)
- As are all of these!
  
- Models can be fit in many ways due to:
  - Insufficient data
  - Noisy data
  - **Model not complex enough**  
(curved boundaries can also separate this data!)



## Ensemble learning II

- Which class (red or blue) should the magenta dot be?
- An obvious classification boundary ...
- But this is just as good (suggesting blue)
- As are all of these!
- Models can be fit in many ways due to:
  - Insufficient data
  - Noisy data
  - **Model not complex enough**  
(curved boundaries can also separate this data!)
- Ensembles have many estimators...  
...to capture this ambiguity



- Ensembles need **diversity**
- Estimators must make **different mistakes**  
i.e. if all make same mistake  $\implies$  ensemble will repeat it
- Increasing estimator diversity at expense of individual performance  $\rightarrow$  better ensemble!  
(up to a limit)

- However, constructing an ensemble would require more data
- Instead of collecting more data...  
    ...fake it from available data

- However, constructing an ensemble would require more data
- Instead of collecting more data...
  - ...fake it from available data
- Given data set of size  $n$ :
  - Create new data set by drawing, with replacement,  $n$  times
  - (there will be repetitions of data-points due to replacement)

- An ensemble technique!
- Short for “Bootstrap AGGREGatING”
- Bagging = Bootstrapping applied to estimator output  
(via optimised parameters)

- An ensemble technique!
- Short for “Bootstrap AGGREGatING”
- Bagging = Bootstrapping applied to estimator output  
(via optimised parameters)
- Algorithm:
  1. Select  $S$ , size of ensemble
  2. Create  $S$  bootstrap draws of original data set
  3. Train estimator on each
  4. Combine outputs of all estimators for each query

## Random subspace method

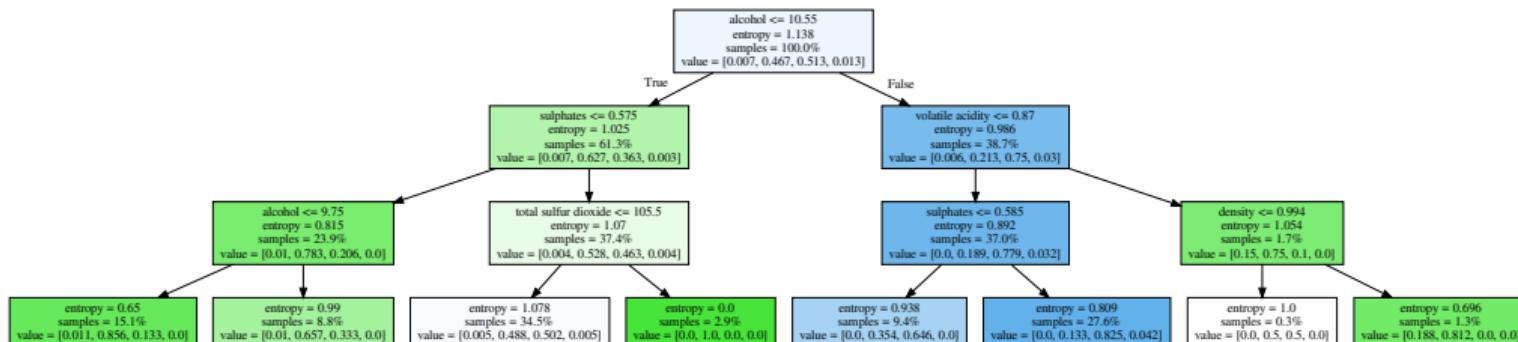
- Another ensemble technique!
- Bootstrap applied to features
  - i.e. fit each estimator with a random subset of features
- For decision tree: new bootstrap for each split

- Random forest = decision trees + bagging + random subspace method
- Algorithm:
  1. Select  $S$ , number of trees (more is better, up to a limit)
  2. Create  $S$  bootstrap draws of original data set
  3. Train decision tree on each, with random subspace method
  4. Combine outputs of all trees for each query
    - Classification : The decision trees vote – ensemble outputs winner
    - Regression : Take the mean/median of all of the estimates

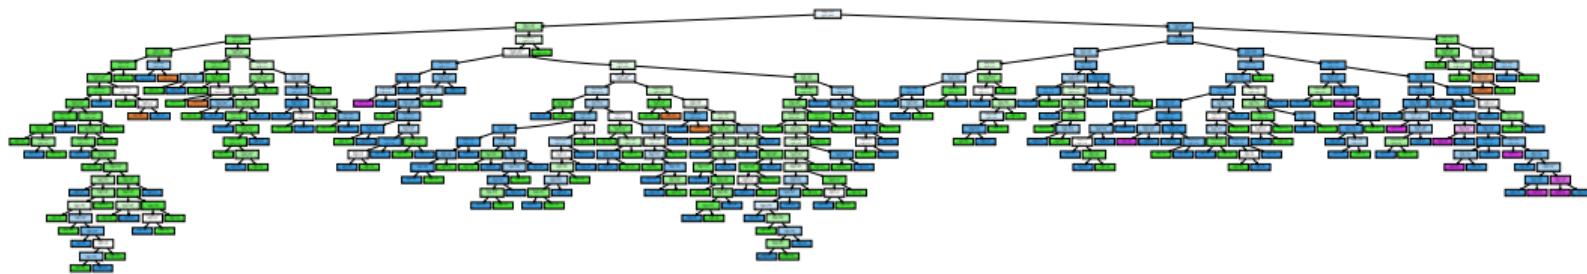
# Explainability

- 1599 exemplars, split: 1199 train, 400 test.
- Input: 11 measurable features:
  - fixed acidity
  - volatile acidity
  - citric acid
  - residual sugar
  - chlorides
  - free sulfur dioxide
  - total sulfur dioxide
  - density
  - pH
  - sulphates
  - alcohol
- Output: 1–10 human rating
  - (reduced to 1–4 here, to fit on screen)
- Can be used for classification or regression!

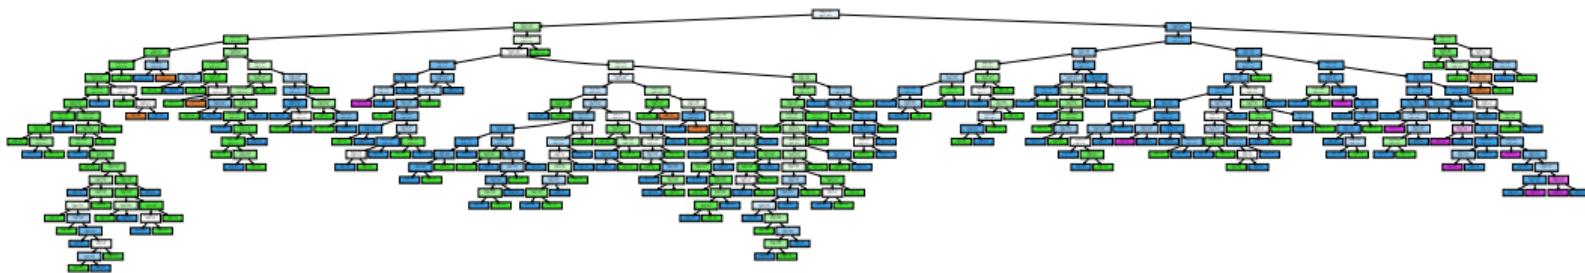
Classification tree, max depth 3:



Accuracy = 69% (and explainable)



- A deeper classification tree:
  - Accuracy = 71% (somewhat less explainable)



- A deeper classification tree:
  - Accuracy = 71% (somewhat less explainable)
- Random forests with 32 trees
  - Classification: Accuracy = 79% (better than 71%)
  - Impossible to visualise/understand!

- “Upgraded” decision trees to
  - Handle more types of inputs & outputs
  - Random forests! (still one of the best)
  - Output probabilities
- Notes:
  - One of the fastest algorithms
  - Many variants, e.g. gradient boosting
- Next week :  
Making sure a ML system is working!

## Further reading

- For more kinds of random forest:  
“Decision Forests for Classification, Regression, Density Estimation, Manifold Learning and Semi-Supervised Learning” by **Criminisi, Shotton and Konukoglu**  
<http://research.microsoft.com/apps/pubs/default.aspx?id=155552>
- For more on ensemble methods:  
“Diversity creation methods: a survey and categorisation”  
by **Brown, Wyatt, Harris and Yao**  
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.421.349&rep=rep1&type=pdf>

Cheese board, CC Worm That Turned, Attribution-Share Alike 4.0 International,  
[https://commons.wikimedia.org/wiki/File:Welsh\\_cheese\\_board.JPG](https://commons.wikimedia.org/wiki/File:Welsh_cheese_board.JPG)

Wine data set from <https://archive.ics.uci.edu/ml/datasets/wine+quality>