

Foundational Machine Learning 1.04: Loss Functions and Performance Measures : Is the ML System Working?

Rohit Babbar
rb2608@bath.ac.uk



- For binary classification, the accuracy of a classifier f on an input-output pair x, y is measured by 0-1 loss :

$$\ell(y, f(x)) = \begin{cases} 1 & \text{if } f(x) \neq y \\ 0 & \text{otherwise} \end{cases}$$

- For binary classification, the accuracy of a classifier f on an input-output pair x, y is measured by 0-1 loss :

$$\ell(y, f(x)) = \begin{cases} 1 & \text{if } f(x) \neq y \\ 0 & \text{otherwise} \end{cases}$$

- On a sample of size n , training/test error of f is given by $R_{emp}(f) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i))$

Terminology

- For binary classification, the accuracy of a classifier f on an input-output pair x, y is measured by 0-1 loss :

$$\ell(y, f(x)) = \begin{cases} 1 & \text{if } f(x) \neq y \\ 0 & \text{otherwise} \end{cases}$$

- On a sample of size n , training/test error of f is given by $R_{emp}(f) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i))$
- However, the Expected (true) loss of f

$$R(f) \triangleq \mathbb{E}_{(x,y) \sim P}(\ell(y, f(x)))$$

The above expectation $\mathbb{E}_{(x,y) \sim P}$ means that it is an expectation (average) that is computed over samples (x, y) drawn from the data distribution P

- For binary classification, the accuracy of a classifier f on an input-output pair x, y is measured by 0-1 loss :

$$\ell(y, f(x)) = \begin{cases} 1 & \text{if } f(x) \neq y \\ 0 & \text{otherwise} \end{cases}$$

- On a sample of size n , training/test error of f is given by $R_{emp}(f) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i))$
- However, the Expected (true) loss of f

$$R(f) \triangleq \mathbb{E}_{(x,y) \sim P}(\ell(y, f(x)))$$

The above expectation $\mathbb{E}_{(x,y) \sim P}$ means that it is an expectation (average) that is computed over samples (x, y) drawn from the data distribution P

- Bayes classifier f_{Bayes} , is defined to be the one which has the least classification error, i.e., $f_{Bayes} = \arg \min_f R(f) := \mathbb{E}_P(\ell(y, f(x)))$

Loss functions in Machine Learning

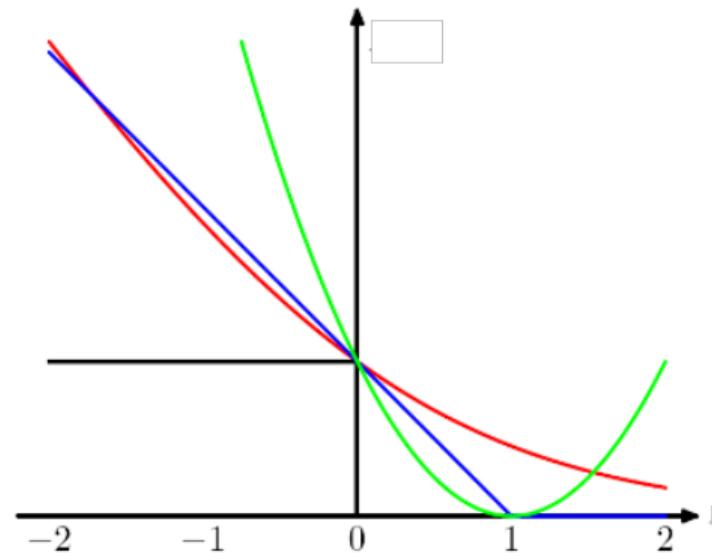


Figure: Horizontal axis is $yf(x)$, vertical is the loss i.e. $\ell(yf(x))$

Convex Upper Bounds on 0-1 loss

- Hinge Loss (in blue) is given by $\max(1 - yf(x), 0)$
- Logistic Loss (in red) is given by $\frac{1}{\log 2} \log(1 + \exp(-yf(x)))$

Bayes Classifier - (1)

Let's say $C_1 = +1$ (Positive class), and $C_2 = -1$ (Negative class) in the figure below. Also $P(\cdot)$ in the text (below) refers to the probability and $p(\cdot)$ in the picture refers to its density.

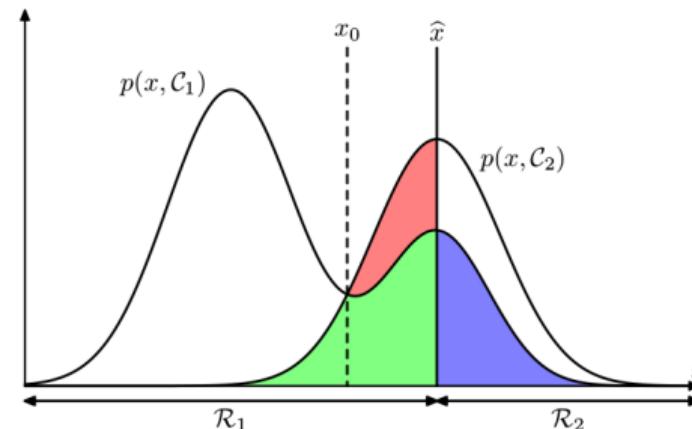
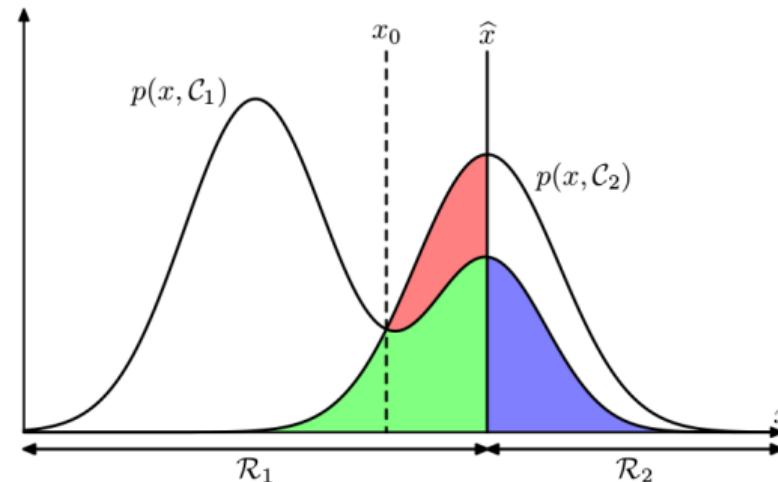


Figure: Depiction of noisy labels (picture from Chris Bishop's book)

- The prediction function of f_{Bayes} is given by

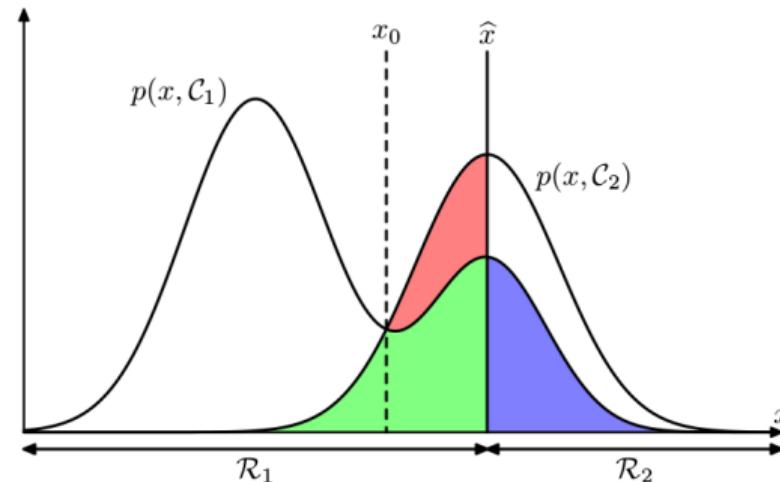
$$f_{Bayes}(x) = \begin{cases} C_1 & \text{if } P(y = C_1 | X = x) \geq 0.5 \\ C_2 & \text{otherwise} \end{cases}$$

Example



- Given a point, say $x = \hat{x}$, how do we compute $P(y = C_1 | X = \hat{x})$ or $P(y = C_2 | X = \hat{x})$?

Example



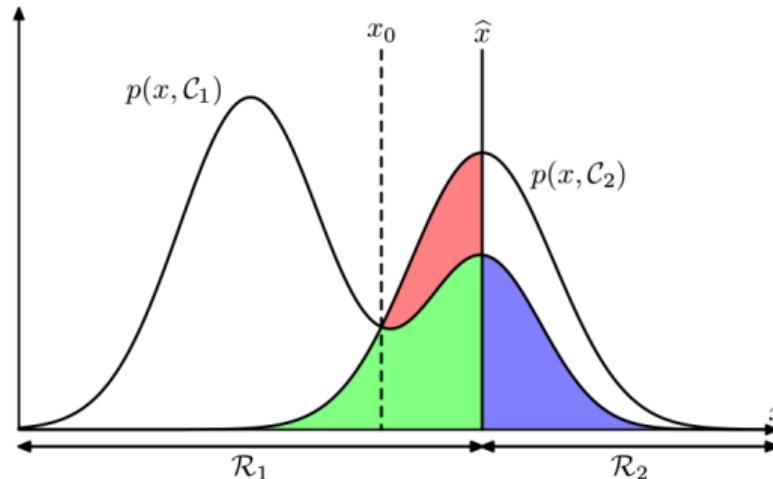
- Given a point, say $x = \hat{x}$, how do we compute $P(y = C_1|X = \hat{x})$ or $P(y = C_2|X = \hat{x})$?

$$P(y = C_1|X = \hat{x}) \stackrel{(1)}{=} \frac{p(X = \hat{x}, Y = C_1)}{p_X(X = \hat{x})} \quad \text{and} \quad P(y = C_2|X = \hat{x}) \stackrel{(1)}{=} \frac{p(X = \hat{x}, Y = C_2)}{p_X(X = \hat{x})}$$

Since the denominator is the same for both, this implies that both probabilities i.e. $P(y = C_i|X = \hat{x})$ for $i = 1, 2$ are proportional to their respective densities at $X = \hat{x}$.

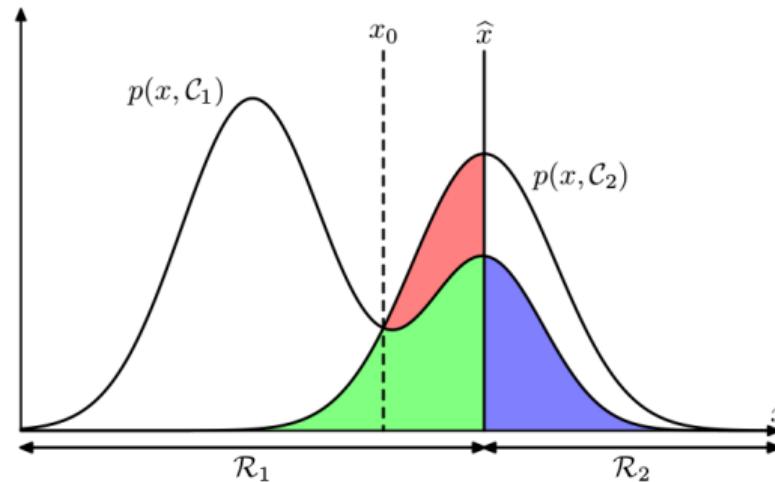
$$\text{Therefore, } P(y = C_1|X = \hat{x}) = \frac{p(X = \hat{x}, Y = C_1)}{p(X = \hat{x}, Y = C_1) + p(X = \hat{x}, Y = C_2)}$$

Example



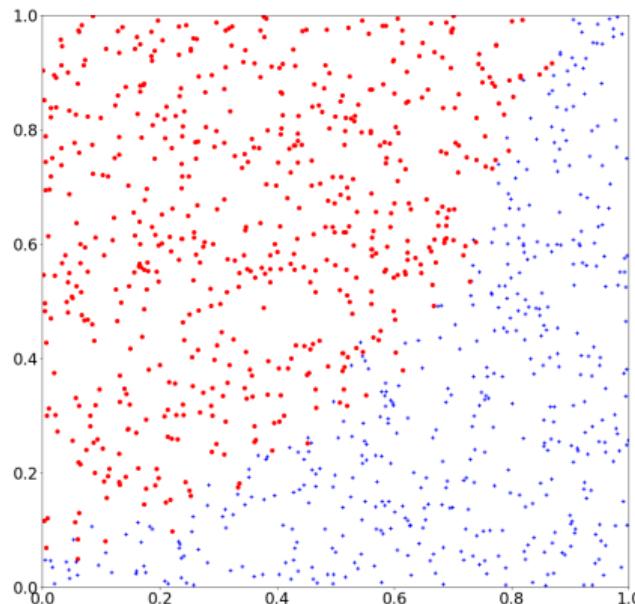
- When the decision boundary/threshold is at $x = \hat{x}$, what kind of errors are signified by the red, green and blue regions?

Example



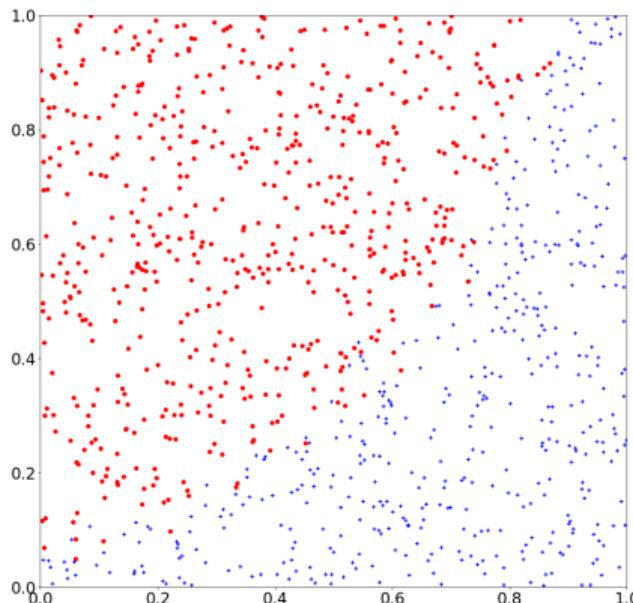
- At what point in the graph $P(y = C_1 | X = x) = 0.5$?

Underfitting & overfitting

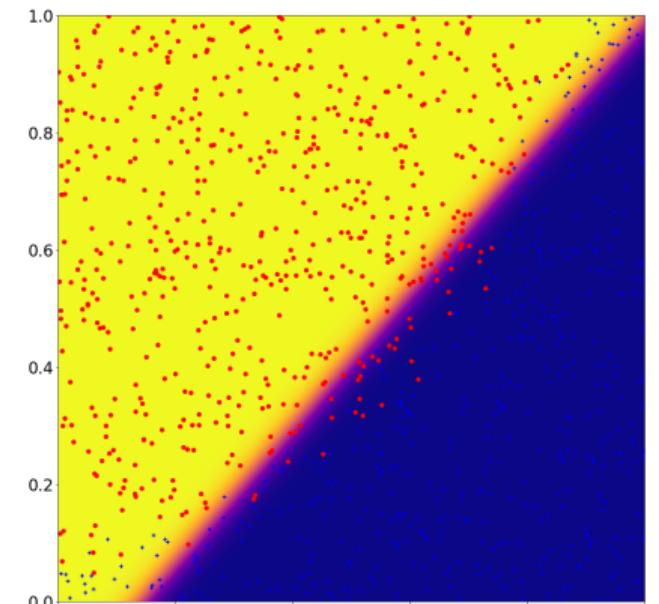


- Curved
- Classes overlap
- How would you divide them?

Underfitting & overfitting

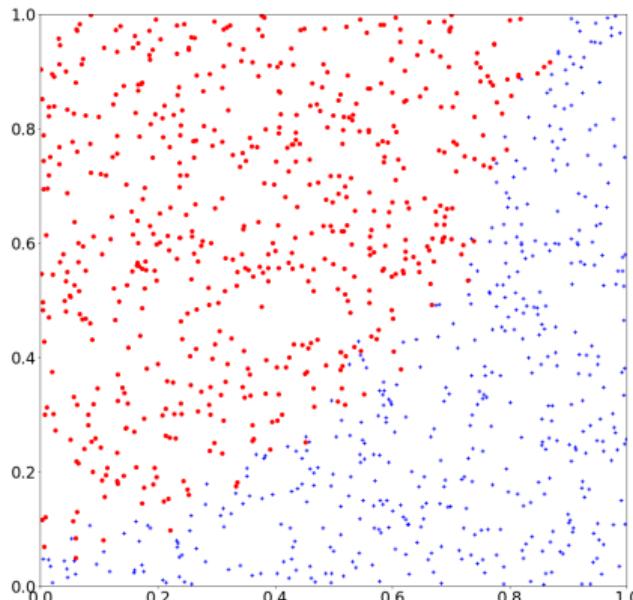


- Curved
- Classes overlap
- How would you divide them?

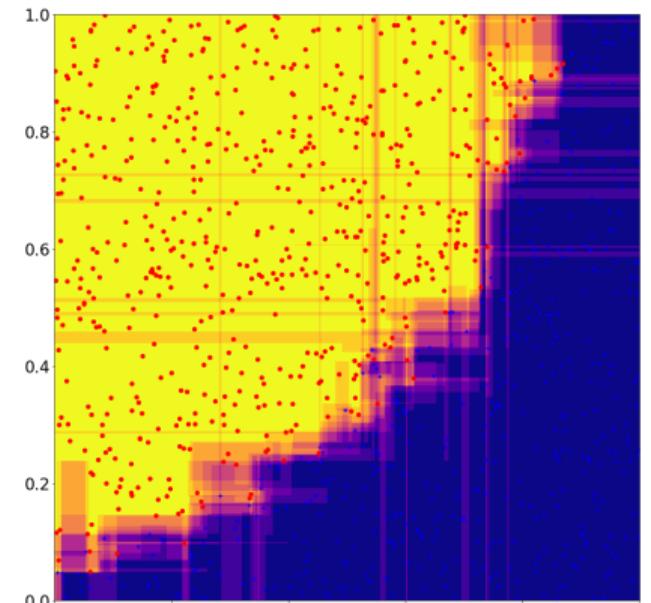


- Underfitting

Underfitting & overfitting



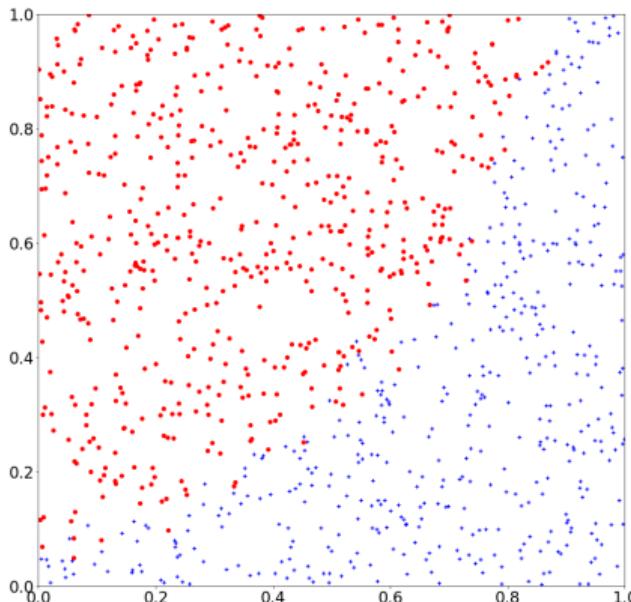
- Curved
- Classes overlap
- How would you divide them?



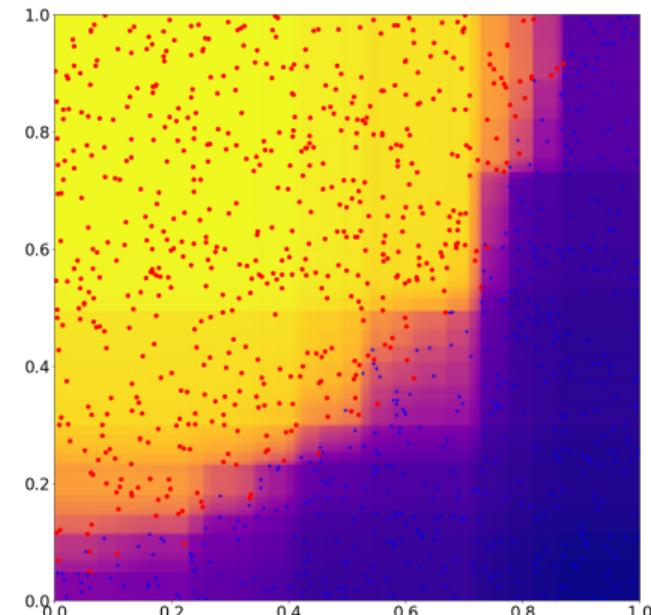
- Overfitting



Underfitting & overfitting

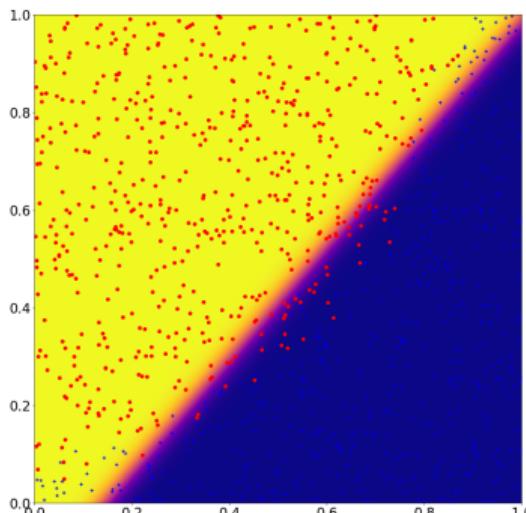


- Curved
- Classes overlap
- How would you divide them?

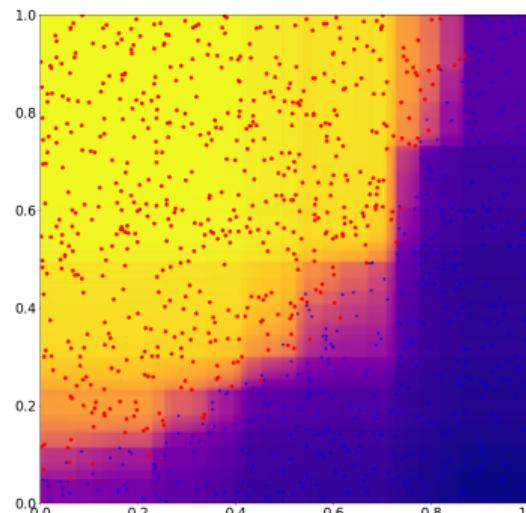


- Balanced

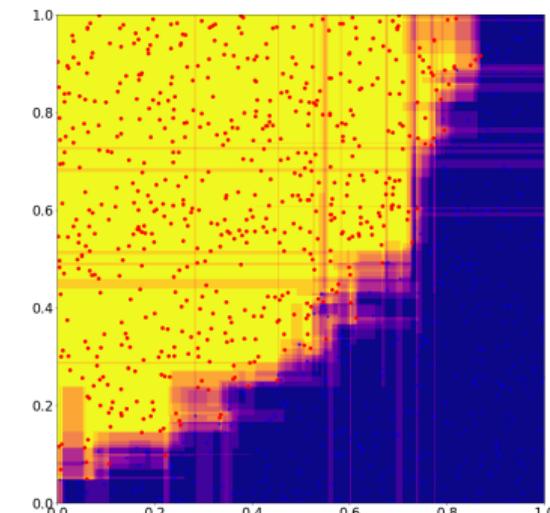
Underfitting & overfitting



- Underfitting
- Logistic regression



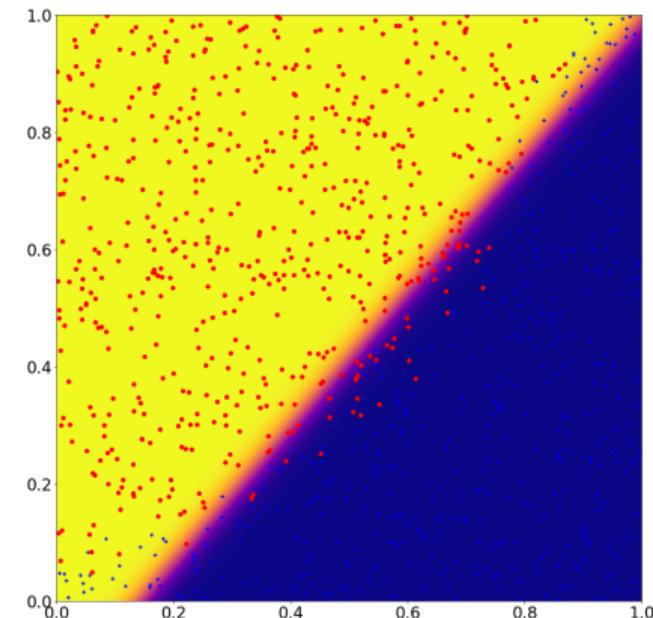
- Balanced
- Tuned random forest
- (scikit learn,
min_impurity_decrease=0.008,
n_estimators=512)



- Overfitting
- Badly tuned random forest
- (scikit learn,
default parameters)

Underfitting causes

- Weak model
- Insufficient data



Overfitting causes

- Powerful model +
- Insufficient **regularisation**
Regularisation \sim smoothing out noise (so model doesn't learn it)
- How to detect?

Train & test set

- Model can't overfit on data it doesn't have!
∴
- Split the data:
 - A **train** set, to fit the model
 - A **test** set, to verify performance

Train & test set

- Model can't overfit on data it doesn't have!
∴
- Split the data:
 - A **train** set, to fit the model
 - A **test** set, to verify performance
- Large gap between train/test accuracy indicates overfitting (usually)

Random Forest	Accuracy	
	Train	Test
Underfitting	79.2%	79.2%
Balanced	97.6%	95.0%
Overfitting	99.6%	94.7%

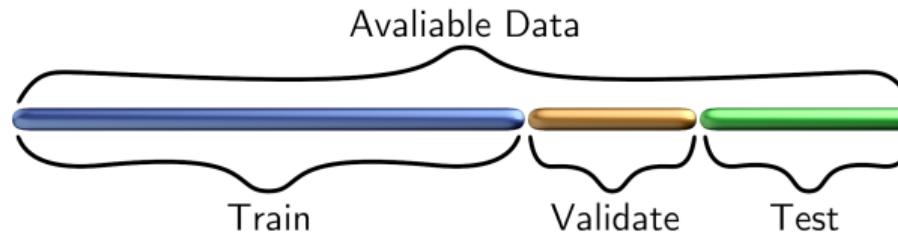
- Parameters → fit to training data
- Hyperparameters → parameter that cannot be fit to training data

- Parameters → fit to training data
- Hyperparameters → parameter that cannot be fit to training data
- Reasons to be a hyperparameter:
 - Avoiding overfitting, e.g. decision tree depth
 - Heavy computation, e.g. ensemble size
 - Hard to optimise

- Can still fit hyperparameters . . .
(manually or by algorithm)
- **...but not to the test set!**
(this mistake can be found in countless research papers)

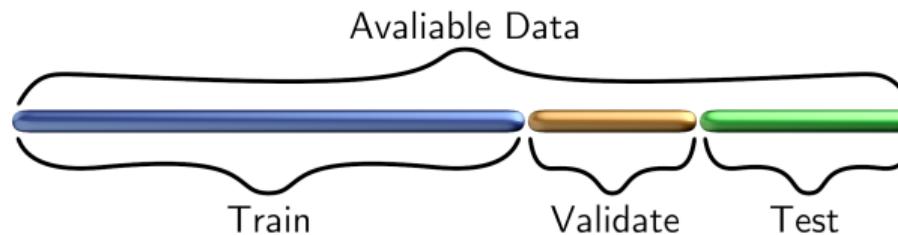
- Can still fit hyperparameters . . .
(manually or by algorithm)
- **...but not to the test set!**
(this mistake can be found in countless research papers)
- Introduce a third set: **validation** set
 - **train** – Give to algorithm
 - **validation** – Objective of hyperparameter optimisation
 - **test** – To report final performance

Measuring performance



- How do we decide on split percentages?
 - Train large → Algorithm performs well
 - Validation large → Hyperparameter optimisation performs well
 - Test large → Accurate performance estimate

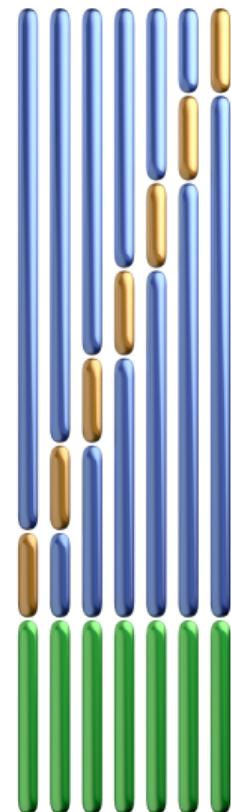
Measuring performance



- How do we decide on split percentages?
 - Train large → Algorithm performs well
 - Validation large → Hyperparameter optimisation performs well
 - Test large → Accurate performance estimate
- Good default: Validation and test small as possible to get reliable estimate, rest on train
- ... but might shrink train due to computational cost
- “small as possible” hard to judge, but 10-15% of data for test and validation each is a good fraction

n-fold

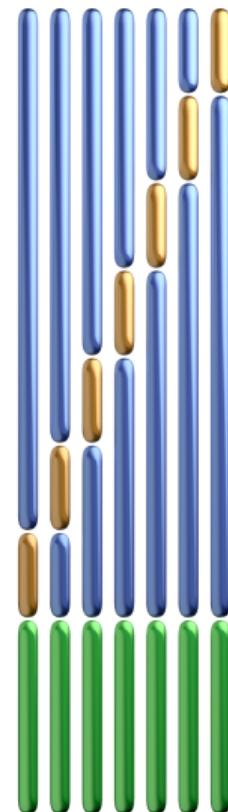
- Validation and test used to make **measurements**
- Can average measurements! (as long as they are independent)



n-fold

- Validation and test used to make **measurements**
- Can average measurements! (as long as they are independent)
- e.g. divide train/validation into 7-fold
 - train: six parts
 - validation: one part

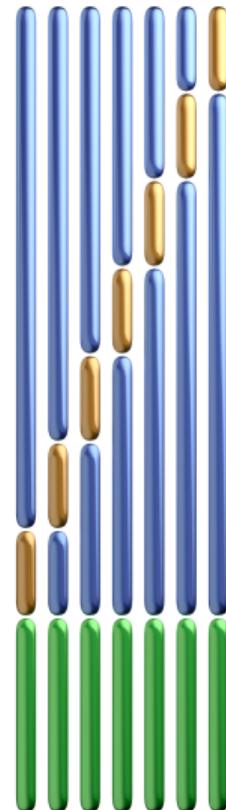
Train for all seven combinations and report average performance on test



- Validation and test used to make **measurements**
- Can average measurements! (as long as they are independent)
- e.g. divide train/validation into 7-fold
 - train: six parts
 - validation: one part

Train for all seven combinations and report average performance on test

- n -fold = $n \times$ slower! Typically $4 \leq n \leq 20$
- General case: All combinations of train/validation/test



- May train algorithm thousands of times!
(hyperparameter tuning and n -fold)
- Choice of n is a trade-off between accuracy / time
- Fast computer/cluster/distributed computation really help!
- Final model: Train on entire data set
(still wise to keep a test set back to sanity check)

Performance?

- What do we actually measure?
(and hence optimise)

- Classification only
- Random forest on breast cancer:

		Actual	
		False	True
Predicted	False	49	6
	True	14	159

Confusion matrices

- Classification only
- Random forest on breast cancer:

		Actual	
		False	True
Predicted	False	49	6
	True	14	159

- On diagonal means correct, off means wrong
- Can see which classes are confused
- An empty row is a problem

Naming the numbers

		Actual	
		False	True
Predicted	False	True Negative (TN)	False Negative (FN)
	True	False Positive (FP)	True Positive (TP)

Naming more numbers

Loads of terms are used :

$$\begin{aligned} & \frac{\text{TP}}{\text{TP} + \text{FN}} \\ & \frac{\text{TN}}{\text{TN} + \text{FP}} \\ & \frac{\text{TP}}{\text{TP} + \text{FP}} \\ & \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}} \\ & \frac{2 \times \text{TP}}{2 \times \text{TP} + \text{FP} + \text{FN}} \end{aligned}$$

sensitivity, **recall**, hit rate, **true positive rate**
specificity, **true negative rate**
precision, positive predictive value
accuracy
F1 score

- Ideally one would like to have high values for both Precision and Recall
- However, between the two, it is typically ease to optimize one metric at the cost of the other
- Optimising F-measure, which is the harmonic mean (i.e. arithmetic mean of the inverses), of Precision and Recall, forces to optimize the smaller one, and hence a better measure

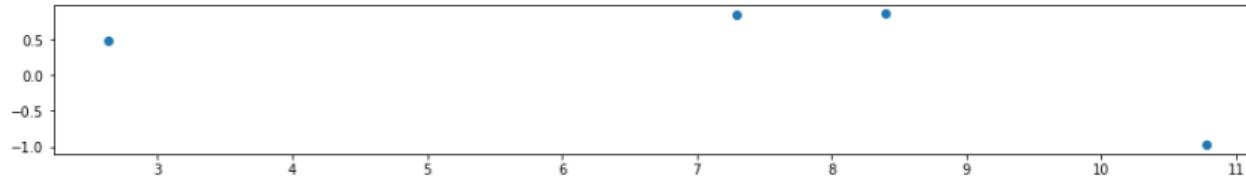
- Imbalanced training set → Makes training difficult
- e.g. credit card fraud $\approx 0.1\%$ of transactions
- 99.9% accuracy by assuming no fraud – meaningless!
- Training : often need to adjust training process (e.g. oversampling)
- Evaluation : F1 score is a better measure
- Balanced accuracy:

$$= \frac{1}{|C|} \sum_{c \in C} \frac{|\{y_i = c \wedge f_\theta(x_i) = c\}|}{|\{y_i = c\}|}$$

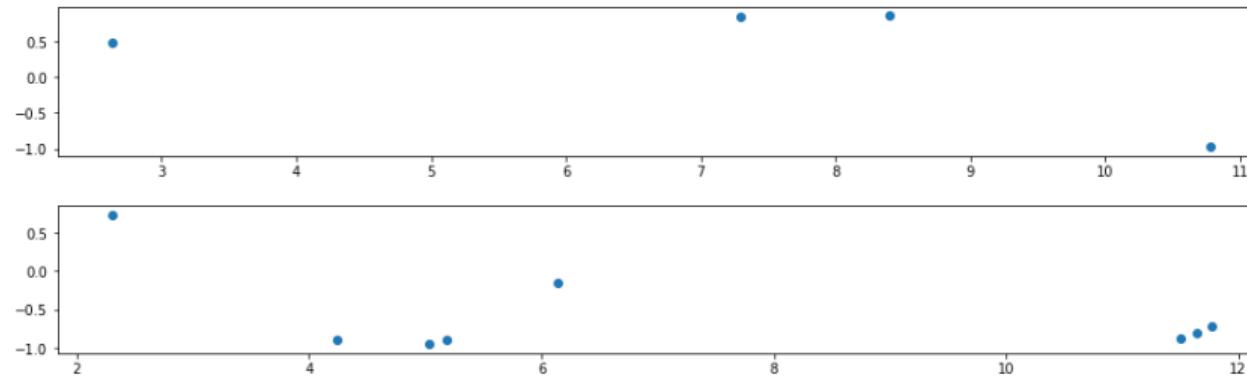
- C = set of classes, of size $|C|$
- (y_i, x_i) = data points
- $f_\theta(\cdot)$ = machine learning model

- Previous are intermediates
- Need a problem specific function of the confusion matrix (for classification)
- Depending on problem might be better to think in terms of:
 - Cost / Loss
 - Gain
 - Error
 - Risk
 - Ranking
- Test complete system!
intermediate proxies can sometimes be misleading

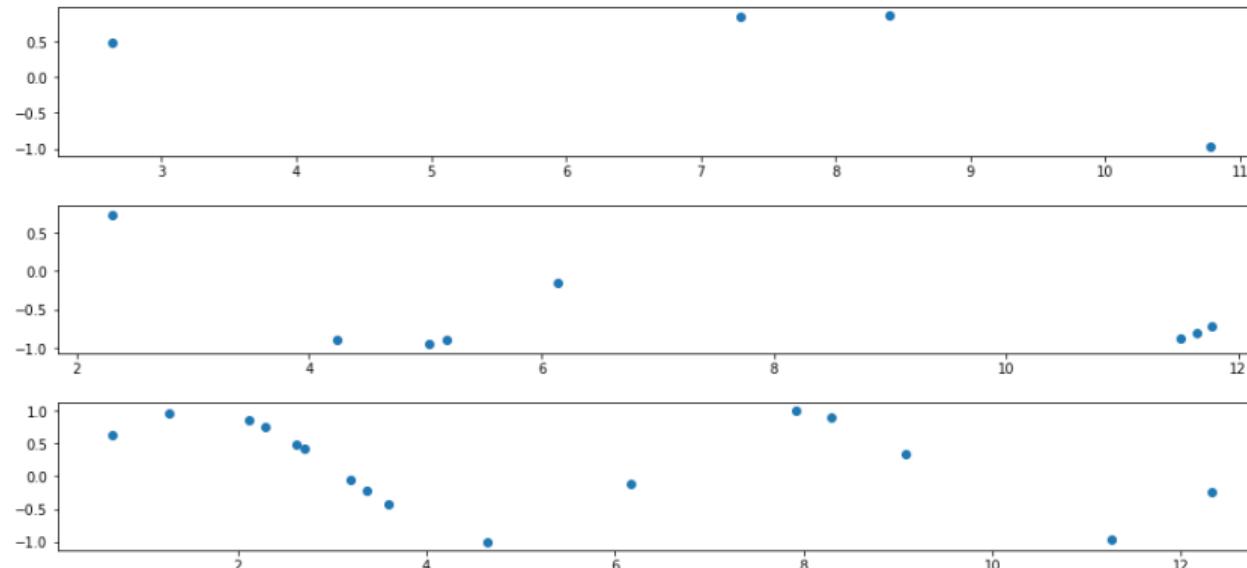
Bad data: Insufficient



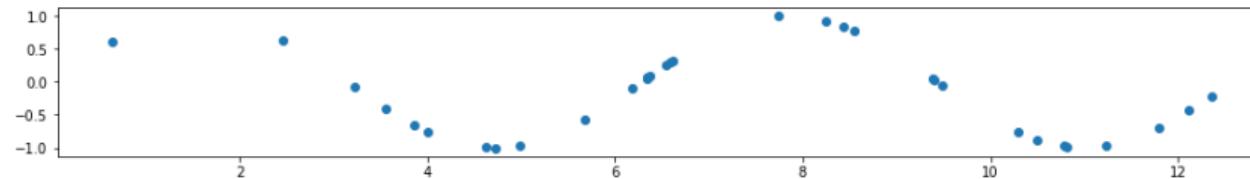
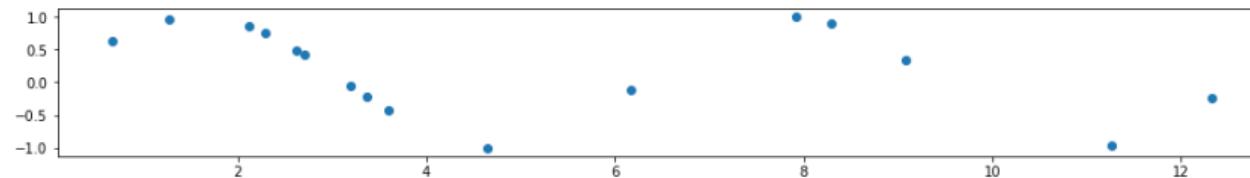
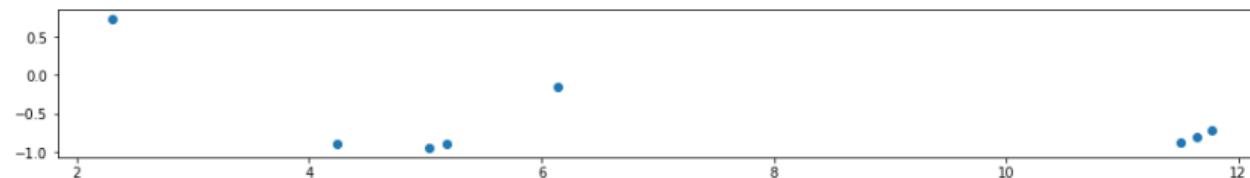
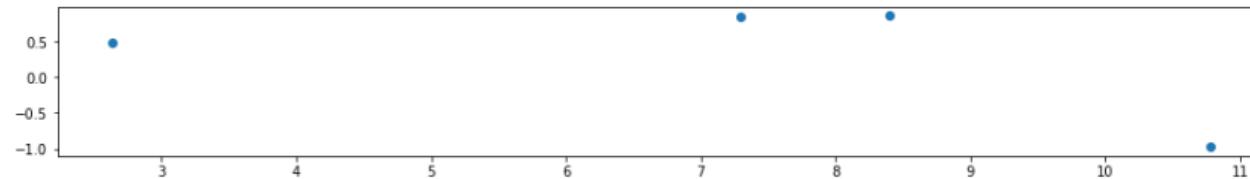
Bad data: Insufficient



Bad data: Insufficient



Bad data: Insufficient



Bad data: Spurious correlation

- 1964: Researcher was spotting M-48 tanks in images
- Got a near perfect score
- Problem:
 - Tank photos taken on a cloudy day
 - Not-tank photos taken on a sunny day
 - ... so it was checking the brightness (b&w so no colour)
- Original paper (probably!): <https://dl.acm.org/citation.cfm?doid=800257.808903>

Bad data: No correlation

- Problem: Estimate when next bus will arrive
- Input:
 - Current height of the fountain
 - Number of purple cars on campus
 - How many bats are in the bat cave
- What's the problem?
- There is nothing to learn – no correlation – it's impossible!

Bad data: Unbalanced

- When you train with 1000 examples of one class and 10 of another
- Classifier can get 99% by always predicting the larger class...
 ... and often does
- Good example: <https://arxiv.org/pdf/1606.08390.pdf>
 - Visual question answering:
 $f(\text{image}, \text{question about image}) \rightarrow \text{answer}$
 - Always giving most common answer ("yes") beat sophisticated approaches!

Bad data: Runtime mismatch

Volvo's driverless cars 'confused' by kangaroos

© 27 June 2017



GETTY IMAGES

| There are more than 20,000 kangaroo strikes each year in Australia

Bad data: Biased data

- Amazon has too many job applications
- ML system sorts best to worst to save time
- Misogynistic
- Trained on past applicants... mostly male
- Spurious correlation as well as bad data
- Article: <https://www.reuters.com/article/us-amazon-com-jobs-automation-insight/amazon-scaps-secret-ai-recruiting-tool-that-showed-bias-against-women-idUSKCN1MK08G>

- Overfitting/underfitting
- Train/test/verification
- Measuring success
- Misinterpretation
- (many kinds of) bad data

- Spotting issues is a skill... takes practise

Further reading

- Blog breaking down why a medical data set is useless: <https://lukeoakdenrayner.wordpress.com/2017/12/18/the-chestxray14-dataset-problems/>
- “Concrete Problems in AI Safety”
by **Amodei, Olah, Steinhardt, Christiano, Schulman and Mane**
<https://arxiv.org/pdf/1606.06565.pdf>
- “How to recognize AI snake oil”
by **Narayanan** (slides)
<https://www.cs.princeton.edu/~arvindn/talks/MIT-STS-AI-snakeoil.pdf>