

Optimisation Basics 3 — Practice Questions

CM52054: Foundational Machine Learning
Practice set with fully worked answers

Part A — Core Concepts

1) Normal equation derivation.

Given the optimisation problem:

$$w^* = \arg \min_w \|Xw - y\|^2,$$

derive the closed-form solution using calculus.

Answer:

$$\nabla_w \|Xw - y\|^2 = 2X^\top Xw - 2X^\top y.$$

Setting the gradient to zero:

$$2X^\top Xw^* - 2X^\top y = 0 \Rightarrow w^* = (X^\top X)^{-1}X^\top y.$$

2) Gradient descent update rule.

Write the update rule for gradient descent applied to the same objective.

Answer:

$$w_{t+1} = w_t - \alpha \nabla f(w_t) = w_t - \alpha(2X^\top Xw_t - 2X^\top y).$$

3) Effect of learning rate.

State what happens if the learning rate α is:

- a) too small;
- b) too large.

Answer:

- a) If α is too small, convergence is slow.
- b) If α is too large, optimisation may oscillate or diverge.

4) Interpretation of the Hessian.

What does the Hessian represent in optimisation?

Answer: The Hessian encodes curvature of the objective function. Large eigenvalues indicate steep curvature requiring small steps; small eigenvalues indicate flatter regions where larger steps are possible.

Part B — Newton's Method

5) Deriving Newton's method.

Using the second-order Taylor expansion of $g(w)$ around w_t , derive the Newton update rule.

Answer: Using Taylor expansion:

$$g(w) \approx g(w_t) + \nabla g(w_t)(w - w_t) + \frac{1}{2}(w - w_t)^\top \nabla^2 g(w_t)(w - w_t).$$

Differentiating and setting to zero:

$$\nabla g(w_t) + \nabla^2 g(w_t)(w^* - w_t) = 0,$$

so

$$w^* = w_t - [\nabla^2 g(w_t)]^{-1} \nabla g(w_t).$$

6) Why one-step convergence for linear regression?

Answer: For linear regression,

$$g(w) = \|Xw - y\|^2, \quad \nabla g(w) = 2X^\top Xw - 2X^\top y, \quad \nabla^2 g(w) = 2X^\top X.$$

Since the Hessian is constant (independent of w),

$$w_{t+1} = w_t - (2X^\top X)^{-1}(2X^\top Xw_t - 2X^\top y) = (X^\top X)^{-1}X^\top y,$$

which is already the optimum, so convergence happens in one step.

7) Why Newton's method is rarely used in large-scale ML.

List three reasons.

Answer:

- Requires computing and inverting the Hessian (computationally expensive).
- Requires second-order derivatives, which may not be available or smooth.
- Only guarantees local convergence and is sensitive to poor initialisation.

8) Gradient descent vs Newton: anisotropic contours.

Explain why gradient descent struggles under anisotropic level sets and how Newton's method resolves this.

Answer: Gradient descent follows the steepest descent direction, causing zig-zag behaviour in elongated (anisotropic) contours. Newton's method rescales the gradient using the inverse Hessian, effectively normalising curvature and directing steps toward the optimum more efficiently.

Part C — More Challenging Questions

9) Newton's method with damping.

In practice, a modified version of Newton's method is often used:

$$w_{t+1} = w_t - \alpha_t [\nabla^2 g(w_t)]^{-1} \nabla g(w_t),$$

where $\alpha_t \in (0, 1]$ is a damping factor. Explain why using $\alpha_t = 1$ may fail even if the Hessian is invertible.

Answer:

Using $\alpha_t = 1$ may fail because the Newton step assumes the second-order Taylor approximation is accurate locally. If w_t is far from the optimum or if the Hessian is not strictly positive definite, the update may move uphill and cause divergence.

10) **Visual interpretation of anisotropy.**

Based on the contour plots shown in the slides 21-23, explain:

- a) Why gradient descent produces a zig-zag trajectory in anisotropic landscapes.
- b) Why Newton's method produces a more direct trajectory.
- c) How the Hessian relates to the geometry of these contours.

Answer:

- a) Gradient descent always moves in the direction of steepest descent, which is nearly orthogonal to the shortest path to the minimum in elongated level sets.
- b) Newton's method rescales gradients using the inverse Hessian, effectively rotating and scaling the step to point along the principal axes of curvature.
- c) The Hessian defines the curvature and orientation of the ellipse-shaped level sets: its eigenvectors give principal directions, and its eigenvalues determine stretching.

11) **Computational trade-offs.**

Compare gradient descent and Newton's method in terms of:

- a) computational cost per iteration,
- b) number of iterations needed for convergence,
- c) storage requirements.

Answer:

- a) Gradient descent requires only first-order derivatives and is cheap per iteration; Newton's method requires computing and inverting the Hessian, which is expensive.
- b) Newton's method converges much faster (often quadratically), sometimes in a single iteration, while gradient descent typically converges linearly.
- c) Gradient descent stores only the gradient and weights, while Newton's method must store and manipulate the full Hessian matrix.