

Logistic Regression — Coding Practice Questions

CM52054: Foundational Machine Learning
Coding practice set with fully worked answers

1) Implement the sigmoid and logit functions.

Write Python functions that:

- a) Implement the sigmoid function $\sigma(z) = 1/(1 + e^{-z})$.
- b) Implement the logit function $\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$.
- c) Test that $\text{logit}(\sigma(z)) \approx z$ for a range of values of z .

Answer:

```
1 import numpy as np
2
3 def sigmoid(z):
4     """
5         Compute the logistic sigmoid function sigma(z) = 1 / (1 + exp(-
6             z)).
7
8         Parameters
9         -----
10        z : float or np.ndarray
11            Input value or array of values.
12
13        Returns
14        -----
15        np.ndarray
16            Sigmoid of the input, with the same shape as 'z'.
17        """
18        # Use the standard definition: sigma(z) = 1 / (1 + exp(-z)).
19        # NumPy will automatically apply exp element-wise if z is an
20        # array.
21        return 1.0 / (1.0 + np.exp(-z))
22
23
24 def logit(p):
25     """
26         Compute the logit function logit(p) = log(p / (1 - p)).
27
28         Parameters
29         -----
29        p : float or np.ndarray
30            Probabilities in the open interval (0, 1).
31
32        Returns
33        -----
34        np.ndarray
35            Logit of the input probabilities.
36        """
```

```

36     # Convert input to a NumPy array for convenience in vectorised
37     # operations.
38     p = np.asarray(p)
39
40     # Basic input check:
41     # We require probabilities to be strictly between 0 and 1 to
42     # avoid log(0).
43     if np.any(p <= 0) or np.any(p >= 1):
44         raise ValueError("All probabilities must be strictly
45                         between 0 and 1.")
46
47
48     # Apply the logit transform element-wise.
49     return np.log(p / (1.0 - p))
50
51
52     # ---- Simple test: logit(sigmoid(z)) \approx z ----
53
54     # Create a range of test values for z.
55     z_values = np.linspace(-5, 5, num=11) # 11 points from -5 to 5
56
57     # Apply sigmoid, then logit.
58     sig_values = sigmoid(z_values)
59     reconstructed_z = logit(sig_values)
60
61     print("Original z values:      ", z_values)
62     print("Reconstructed z values:", reconstructed_z)
63     print("Difference:           ", reconstructed_z - z_values)

```

Listing 1: Sigmoid and logit implementations with a simple test

Explanation:

- The `sigmoid` function directly implements $1/(1 + e^{-z})$.
- The `logit` function applies $\log(p/(1 - p))$ and checks that p lies strictly in $(0, 1)$ to avoid numerical issues.
- The small test checks that $\text{logit}(\sigma(z))$ approximately recovers z , up to floating point rounding.

2) Implement logistic loss and gradient.

Consider the logistic regression model with

$$p_i = \sigma(w^\top x_i), \quad \ell(w) = -\frac{1}{N} \sum_{i=1}^N [y_i \log p_i + (1 - y_i) \log(1 - p_i)].$$

Assume:

- $X \in \mathbb{R}^{N \times M}$ has rows x_i^\top ,
- $y \in \{0, 1\}^N$,
- $w \in \mathbb{R}^M$.

- Implement a function that returns the average logistic loss $\ell(w)$.
- Implement the gradient $\nabla_w \ell(w) = -\frac{1}{N} \sum_{i=1}^N (y_i - p_i)x_i$.

Answer:

```

1 def logistic_loss_and_grad(w, X, y):
2     """
3         Compute the average logistic loss (binary cross-entropy) and
4             its gradient.
5
6         We assume a logistic regression model:
7             p_i = sigmoid(w^T x_i)
8             and loss:
9                 ell(w) = -1/N * sum_i [ y_i log(p_i) + (1 - y_i) log(1 -
10                p_i) ].
11
12     Parameters
13     -----
14         w : np.ndarray of shape (M,)
15             Parameter vector of the model.
16         X : np.ndarray of shape (N, M)
17             Design matrix; each row is a sample x_i.
18         y : np.ndarray of shape (N,)
19             Binary labels in {0, 1}.
20
21     Returns
22     -----
23         loss : float
24             The average logistic loss over the N samples.
25         grad : np.ndarray of shape (M,)
26             Gradient of the loss with respect to w.
27
28     """
29     # Number of samples.
30     N = X.shape[0]
31
32     # Compute linear scores z_i = w^T x_i for all samples at once.
33     # X has shape (N, M), w has shape (M,), so X @ w -> (N,).
34     z = X @ w
35
36     # Convert scores into probabilities p_i = sigma(z_i).
37     p = sigmoid(z)    # shape (N,)
38
39     # Clip probabilities to avoid log(0) when computing the loss.
40     eps = 1e-12 # 1*10^{-12}
41     p_clipped = np.clip(p, eps, 1.0 - eps)
42
43     # Compute average binary cross-entropy:
44     #     loss = -1/N * sum_i [ y_i log(p_i) + (1-y_i) log(1-p_i) ].
45     loss = -np.mean(y * np.log(p_clipped) + (1 - y) * np.log(1 -
46             p_clipped))
47
48     # Compute the gradient:
49     #     grad = -1/N * sum_i (y_i - p_i) x_i.
50     # First compute the vector of differences (y_i - p_i).
51     diff = y - p    # shape (N,)
52
53     # Then accumulate using X^T @ diff:
54     #     X^T has shape (M, N); diff has shape (N,).
55     #     Result is shape (M,), as desired.
56     grad = -(1.0 / N) * (X.T @ diff)
57
58     return loss, grad

```

Listing 2: Logistic loss and gradient (vectorised)

Explanation:

- We compute the vector of scores $\mathbf{z} = \mathbf{X}\mathbf{w}$ and then apply the sigmoid to get probabilities p_i .
- Probabilities are clipped to avoid numerical issues when taking logs.
- The loss is implemented using vectorised operations and averaged over all samples.
- For the gradient, we build the vector $(y_i - p_i)$ and accumulate $\sum_i(y_i - p_i)x_i$ via $\mathbf{X}^\top \mathbf{diff}$.

3) Train logistic regression with batch gradient descent.

Write a function that:

- Takes (\mathbf{X}, \mathbf{y}) , an initial weight vector \mathbf{w}_{init} , a learning rate α , and a number of iterations `num_iters`.
- Performs batch gradient descent on the logistic loss:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \alpha \nabla_{\mathbf{w}} \ell(\mathbf{w}_t).$$

- Returns the final weights and the history of loss values.

Reuse `logistic_loss_and_grad` from the previous question.

Answer:

```

1 def train_logistic_regression(X, y, w_init, alpha=0.1, num_iters
2     =1000):
3     """
4         Train logistic regression via batch gradient descent.
5
6         Parameters
7         -----
8             X : np.ndarray of shape (N, M)
9                 Training data matrix.
10            y : np.ndarray of shape (N,)
11                Binary labels in {0, 1}.
12            w_init : np.ndarray of shape (M,)
13                Initial parameter vector.
14            alpha : float, optional
15                Learning rate (step size) for gradient descent.
16            num_iters : int, optional
17                Number of gradient descent iterations.
18
19         Returns
20         -----
21             w : np.ndarray of shape (M,)
22                 Final parameter vector after training.
23             loss_history : list of float
24                 List of loss values recorded at each iteration.
25             """
26             # Copy initial weights so we don't modify the caller's array.
27             w = w_init.copy()
28
29             # Store loss values to inspect convergence behaviour afterwards
30             .

```

```

29     loss_history = []
30
31     for t in range(num_iters):
32         # Compute current loss and gradient for the current
33         # parameter vector.
34         loss, grad = logistic_loss_and_grad(w, X, y)
35
36         # Record loss.
37         loss_history.append(loss)
38
39         # Gradient descent update (we subtract because we minimise
40         # the loss).
41         w = w - alpha * grad
42
43         # Optional: print progress every fixed number of iterations
44         .
45         if (t + 1) % 100 == 0:
46             print(f"Iteration {t+1:4d} / {num_iters}, loss = {loss
47                  :.4f}")
48
49     return w, loss_history

```

Listing 3: Batch gradient descent for logistic regression

Explanation:

- The loop runs a fixed number of iterations of gradient descent.
- At each step, we compute both the current loss and its gradient.
- The update rule $w \leftarrow w - \alpha \nabla \ell(w)$ moves us in the direction of steepest descent of the loss.
- The list `loss_history` can be used to plot or check convergence.

4) Prediction and accuracy for logistic regression.

Implement the following functions:

- `predict_proba(X, w)` returning predicted probabilities $p(y = 1 | x_i)$.
- `predict_labels(X, w, threshold)` returning predicted labels in $\{0, 1\}$.
- `accuracy(y_true, y_pred)` computing classification accuracy.

Answer:

```

1 def predict_proba(X, w):
2     """
3         Predict class-1 probabilities p(y=1 | x) for all samples in X.
4
5     Parameters
6     -----
7     X : np.ndarray of shape (N, M)
8         Data matrix.
9     w : np.ndarray of shape (M,)
10        Parameter vector.
11
12    Returns
13    -----
14    p : np.ndarray of shape (N,)
15        Predicted probabilities for class 1.

```

```

16     """
17     # Compute the linear scores  $z = Xw$  for all samples.
18     z = X @ w
19
20     # Convert scores to probabilities using the sigmoid function.
21     p = sigmoid(z)
22     return p
23
24
25 def predict_labels(X, w, threshold=0.5):
26     """
27     Predict binary labels in {0, 1} for all samples in X using
28     logistic regression.
29
30     Parameters
31     -----
32     X : np.ndarray of shape (N, M)
33         Data matrix.
34     w : np.ndarray of shape (M,)
35         Parameter vector.
36     threshold : float, optional
37         Decision threshold on the predicted probability for class
38         1.
39         If  $p \geq \text{threshold}$ , predict 1; otherwise 0.
40
41     Returns
42     -----
43     y_pred : np.ndarray of shape (N,)
44         Predicted labels in {0, 1}.
45     """
46
47     # Get predicted probabilities for each sample.
48     p = predict_proba(X, w)
49
50
51     # Compare each probability to the threshold to get class labels
52     .
53     y_pred = np.where(p >= threshold, 1, 0)
54     return y_pred
55
56
57 def accuracy(y_true, y_pred):
58     """
59     Compute classification accuracy: fraction of correct
60     predictions.
61
62     Parameters
63     -----
64     y_true : np.ndarray of shape (N,)
65         True labels.
66     y_pred : np.ndarray of shape (N,)
67         Predicted labels.
68
69     Returns
70     -----
71     acc : float
72         Accuracy in [0, 1].
73     """
74
75     # Convert inputs to NumPy arrays for safety.
76     y_true = np.asarray(y_true)

```

```

70     y_pred = np.asarray(y_pred)
71
72     # Accuracy is simply the fraction of positions where y_true ==
73     # y_pred.
73     return np.mean(y_true == y_pred)

```

Listing 4: Prediction utilities: probabilities, labels, accuracy

Explanation:

- `predict_proba` just performs the forward pass of the model.
- `predict_labels` thresholds probabilities at a user-specified value (default 0.5).
- `accuracy` calculates the fraction of correct predictions.

5) End-to-end training on synthetic data.

Write a small script that:

- Generates synthetic 2D data with a bias term, using a known “true” parameter vector w_{true} .
- Samples labels from Bernoulli probabilities defined by this w_{true} .
- Trains logistic regression with `train_logistic_regression`.
- Prints the learned weights and the training accuracy.

Answer:

```

1 import numpy as np
2
3 # 1) Generate synthetic data.
4 np.random.seed(42)    # Fix seed for reproducibility.
5
6 N = 200   # Number of samples.
7
8 # True parameters for synthetic data, including bias term:
9 # w_true = [bias, w1, w2].
10 w_true = np.array([-0.5, 2.0, -1.0])
11
12 # Generate 2D features x1, x2 from a standard normal distribution.
13 X_raw = np.random.randn(N, 2)  # Shape (N, 2).
14
15 # Add a column of ones to encode the bias term as a feature.
16 # Resulting X has shape (N, 3): [1, x1, x2].
17 X = np.hstack([np.ones((N, 1)), X_raw])
18
19 # Compute linear scores and probabilities according to w_true.
20 z_true = X @ w_true
21 p_true = sigmoid(z_true)
22
23 # Sample labels y from Bernoulli(p_true).
24 # For each sample i, y_i ~ Bernoulli(p_true[i]).
25 y = (np.random.rand(N) < p_true).astype(int)
26
27 print("First 5 synthetic labels:", y[:5])
28
29 # 2) Train logistic regression via gradient descent.
30
31 # Initial weights, e.g., all zeros.

```

```

32 w_init = np.zeros(X.shape[1])
33
34 # Hyperparameters for training.
35 alpha = 0.1      # Learning rate.
36 num_iters = 1000 # Number of gradient descent steps.
37
38 w_learned, loss_history = train_logistic_regression(
39     X, y, w_init, alpha=alpha, num_iters=num_iters
40 )
41
42 print("\nLearned weights:", w_learned)
43 print("True weights:   ", w_true)
44
45 # 3) Evaluate accuracy on the training set.
46
47 y_pred = predict_labels(X, w_learned, threshold=0.5)
48 train_acc = accuracy(y, y_pred)
49
50 print(f"\nTraining accuracy: {train_acc * 100:.2f}%")

```

Listing 5: End-to-end logistic regression on synthetic data

Explanation:

- A true parameter vector w_{true} defines the ground-truth decision boundary.
- 2D features are drawn from a Gaussian, and a bias column of ones is added.
- Probabilities are computed via $\sigma(Xw_{\text{true}})$, and labels are sampled from the corresponding Bernoulli distributions.
- Logistic regression is trained starting from zero weights. The learned weights are compared to the true ones, and training accuracy is reported.

6) Optional: Visualise the learned decision boundary.

Extend the script from the previous question to:

- Plot the 2D data points in the (x_1, x_2) plane, coloured by class label.
- Overlay the decision boundary given by $w_0 + w_1x_1 + w_2x_2 = 0$ for the learned parameter vector w .

Answer:

```

1 import matplotlib.pyplot as plt
2
3 def plot_decision_boundary(X_raw, y, w):
4     """
5         Plot 2D data points and the decision boundary for a logistic
6         regression model.
7
8         Parameters
9         -----
10        X_raw : np.ndarray of shape (N, 2)
11            The original 2D features (without the bias column).
12        y : np.ndarray of shape (N,)
13            Binary labels in {0, 1}.
14        w : np.ndarray of shape (3,)
15            Learned parameters [bias, w1, w2].
16    """

```

```

16     # Build masks for each class to plot them differently.
17     class0 = (y == 0)
18     class1 = (y == 1)
19
20     # Create a new figure for plotting.
21     plt.figure(figsize=(6, 5))
22
23     # Scatter plot for class 0 samples.
24     plt.scatter(
25         X_raw[class0, 0], X_raw[class0, 1],
26         marker='o', alpha=0.7, label='Class 0',
27     )
28
29     # Scatter plot for class 1 samples.
30     plt.scatter(
31         X_raw[class1, 0], X_raw[class1, 1],
32         marker='s', alpha=0.7, label='Class 1',
33     )
34
35     # Extract parameters: w = [b, w1, w2].
36     b, w1, w2 = w
37
38     # Decision boundary satisfies b + w1*x1 + w2*x2 = 0.
39     # Solve for x2 in terms of x1: x2 = -(b + w1*x1) / w2.
40     x1_min, x1_max = X_raw[:, 0].min() - 1.0, X_raw[:, 0].max() +
41         1.0
41     x1_vals = np.linspace(x1_min, x1_max, 100)
42
43     # If w2 is very close to zero, the boundary is almost vertical.
44     if np.abs(w2) < 1e-8:
45         x1_boundary = -b / w1
46         plt.axvline(x=x1_boundary, linestyle='--', label='Decision
47             boundary')
48     else:
49         x2_vals = -(b + w1 * x1_vals) / w2
50         plt.plot(x1_vals, x2_vals, 'k--', label='Decision boundary',
51             )
52
53     plt.xlabel('x1')
54     plt.ylabel('x2')
55     plt.legend()
56     plt.title('Logistic Regression Decision Boundary')
57     plt.grid(True)
58     plt.show()
59
60     # Usage with the synthetic data and learned parameters from Q5:
61     plot_decision_boundary(X_raw, y, w_learned)

```

Listing 6: Plotting the 2D data and learned decision boundary

Explanation:

- Data points for the two classes are plotted with different markers, making the separation visible.
- The decision boundary is obtained from the equation $w_0 + w_1x_1 + w_2x_2 = 0$.
- We either plot x_2 as a function of x_1 , or draw a vertical line if the boundary is nearly vertical.